

Llamadas al sistema

Programando llamadas al sistema en forma manual

Hello World!!

- Para programar un hello world necesitamos hacer 2 llamadas al sistema:
 - Escribir en pantalla el mensaje
 - Terminar el proceso
- Por ello tendremos que hacer uso de las siguientes llamadas al sistema:
 - write (man 2 write)
 - exit (man exit)

Hello World en x86-32bit

- En x86-32bit las system calls tienen los siguientes números:
 - write → syscall número 4
 - exit → syscall número 1
- Linux-2.6.34.14/arch/x86/include/asm/unistd_32.h

```
/*  
 * This file contains the system call numbers.  
 */  
  
#define NR_restart_syscall 0  
#define NR_exit 1  
#define NR_fork 2  
#define NR_read 3  
#define NR_write 4  
#define NR_open 5  
#define NR_close 6  
#define NR_waitpid 7  
#define NR_creat 8  
#define NR_link 9  
#define NR_unlink 10  
#define NR_execve 11  
#define NR_chdir 12  
#define NR_time 13
```


Hello World en x86-32bit (cont)

- Los manuales de las system calls permiten saber cuales son los parámetros

NAME

write - write to a file descriptor

SYNOPSIS

```
#include <unistd.h>
```

```
ssize_t write(int fd, const void *buf, size_t count);
```

DESCRIPTION

`write()` writes up to count bytes from the buffer pointed buf to the file referred to by the file descriptor fd.

NAME

exit - cause normal process termination

SYNOPSIS

```
#include <stdlib.h>
```

```
void exit(int status);
```

DESCRIPTION

The `exit()` function causes normal process termination and the value of status & 0377 is returned to the parent (see `wait(2)`).

NAME

write - write to a file descriptor

SYNOPSIS

```
#include <unistd.h>
```

```
ssize_t write(int fd, const void *buf, size_t count);
```

DESCRIPTION

`write()` writes up to count bytes from the buffer pointed buf to the file referred to by the file descriptor fd.

```
_start:

; sys_write(stdout, message, length)

mov eax, 4      ; sys_write syscall
mov ebx, 1      ; stdout
mov ecx, message ; message address
mov edx, length  ; message string length
int 80h

; sys_exit(return_code)

mov eax, 1      ; sys_exit syscall
mov ebx, 0      ; return 0 (success)
int 80h

section .data
message: db 'Hello, world!',0x0A ; message and newline
```

NAME

exit - cause normal process termination

SYNOPSIS

```
#include <stdlib.h>
```

```
void exit(int status);
```

DESCRIPTION

The `exit()` function causes normal process termination and the value of status & 0377 is returned to the parent (see `wait(2)`).

Hello World en x86-64bit

- En x86-64bit las system calls tienen los siguientes números:
 - write → syscall número 1
 - exit → syscall número 60
- Linux-2.6.34.14/arch/x86/include/asm/unistd_64.h

```
/*
 * This file contains the system call numbers.
 *
 * Note: holes are not allowed.
 */

/* at least 8 syscall per cacheline */
#define __NR_read 0
__SYSCALL(__NR_read, sys_read)
#define __NR_write 1
__SYSCALL(__NR_write, sys_write)
#define __NR_open 2
__SYSCALL(__NR_open, sys_open)
#define __NR_close 3
__SYSCALL(__NR_close, sys_close)
#define __NR_stat 4
```

```
#define __NR_vfork 58
__SYSCALL(__NR_vfork, stub_vfork)
#define __NR_execve 59
__SYSCALL(__NR_execve, stub_execve)
#define __NR_exit 60
__SYSCALL(__NR_exit, sys_exit)
#define __NR_wait4 61
__SYSCALL(__NR_wait4, sys_wait4)
#define __NR_kill 62
__SYSCALL(__NR_kill, sys_kill)
#define __NR_uname 63
```

Hello World en x86-64bit (cont)

- Se usan los mismos manuales de las system call que para el caso anterior
- Lo que cambia es el número de system call, el cual está en el kernel, no en el manual de la misma.
- Los procesadores X86-64 usan un esquema de registros diferentes.
- Se usa la instrucción **syscall** en lugar de la instrucción **int 80h**

NAME

write - write to a file descriptor

SYNOPSIS

```
#include <unistd.h>
```

```
ssize_t write(int fd, const void *buf, size_t count);
```

DESCRIPTION

`write()` writes up to count bytes from the buffer pointed buf to the file referred to by the file descriptor fd.

```
; sys_write(stdout, message, length)

mov    rax, 1      ; sys_write
mov    rdi, 1      ; stdout
mov    rsi, message ; message address
mov    rdx, length ; message string length
syscall

; sys_exit(return_code)

mov    rax, 60     ; sys_exit
mov    rdi, 0      ; return 0 (success)
syscall

section .data
message: db 'Hello, world!', 0x0A ; message and newline
length:  equ $-message           ; NASM definition pseudo-instruction
```

NAME

exit - cause normal process termination

SYNOPSIS

```
#include <stdlib.h>
```

```
void exit(int status);
```

DESCRIPTION

The `exit()` function causes normal process termination and the value of status & 0377 is returned to the parent (see `wait(2)`).

Referencias

Como programar un “hello world” en x86 32bit y 64bit

- <http://blog.markloiseau.com/2012/05/64-bit-hello-world-in-linux-assembly-na-sm/>

Alternativas para pasar parametros a las system calls

- <http://www.int80h.org/bsdasm/#system-calls>